

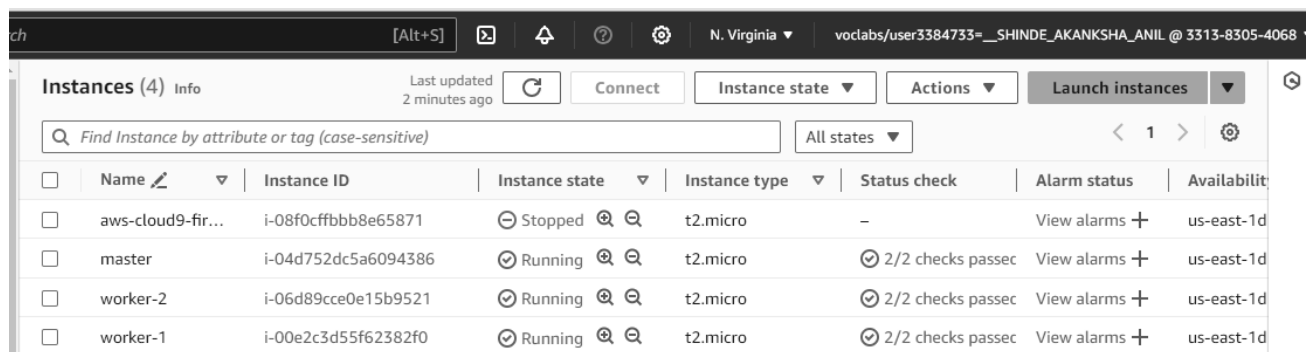
Experiment: 3

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Steps:

1. Create 3 EC2 Ubuntu Instances on AWS.

(Name 1 as Master, the other 2 as worker-1 and worker-2)



	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability
<input type="checkbox"/>	aws-cloud9-fir...	i-08f0cffbb8e65871	Stopped	t2.micro	-	View alarms +	us-east-1d
<input type="checkbox"/>	master	i-04d752dc5a6094386	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1d
<input type="checkbox"/>	worker-2	i-06d89cce0e15b9521	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1d
<input type="checkbox"/>	worker-1	i-00e2c3d55f62382f0	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1d

2. Edit the Security Group Inbound Rules to allow SSH and do it for all the three machines.

```
Akanksha Shinde@AkankshaShinde MINGW64 ~/downloads (master)
$ ssh -i "server.pem" ec2-user@ec2-54-174-206-93.compute-1.amazonaws.com
The authenticity of host 'ec2-54-174-206-93.compute-1.amazonaws.com (54.174.206.93)' can't be established.
ED25519 key fingerprint is SHA256:T+tsGyI15gAvUvjeAZ7GjDIWXH0aI4EPF5g5oICrkoQ.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-54-174-206-93.compute-1.amazonaws.com' (ED25519)
to the list of known hosts.

_#_
~\#####_ Amazon Linux 2023
~~\#####\
~~\###|
~~\#/_ https://aws.amazon.com/linux/amazon-linux-2023
```

3. From now on, until mentioned, perform these steps on all 3 machines. Install Docker for all the 3 machines

```
[root@ip-172-31-90-172 ec2-user]# yum install docker -y
Last metadata expiration check: 0:21:16 ago on Fri Aug 30 04:01:12 2024.
Dependencies resolved.
```

Package	Architecture	Version
Installing:		
docker	x86_64	25.0.6-1.amzn2023.0.1
Installing dependencies:		
containerd	x86_64	1.7.20-1.amzn2023.0.1
iptables-libs	x86_64	1.8.8-3.amzn2023.0.2
iptables-nft	x86_64	1.8.8-3.amzn2023.0.2
libcgroup	x86_64	3.0-1.amzn2023.0.1
libnetfilter_conntrack	x86_64	1.0.8-2.amzn2023.0.2
libnftnl	x86_64	1.0.1-19.amzn2023.0.2
libnftnl	x86_64	1.2.2-2.amzn2023.0.2
pigz	x86_64	2.5-1.amzn2023.0.3

Start the docker by running the command `systemctl start docker` in the terminal of all the ec2 instance.

```
Complete!
[root@ip-172-31-82-133 ec2-user]# systemctl start docker
[root@ip-172-31-82-133 ec2-user]#
```

4. Install the kubernetes on all 3 machines by searching for kubeadm and click on install kubernetes.

Select the red hat based distribution. This process will automatically disable SELinux before configuring kubelet so no need to run it separately in terminal.

version.

Debian-based distributions

Red Hat-based distributions

Without a package manager

1. Set SELinux to `permissive` mode:

These instructions are for Kubernetes 1.31.

```
# Set SELinux in permissive mode (effectively disabling it)
sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/
```

Copy the below script, to install kubernetes we need a kubernetes repo so this script helps us in getting that and paste it in the terminal.

```
# This overwrites any existing configuration in /etc/yum.repos.d/kubern
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
```

```

Installed:
  containerd-1.7.20-1.amzn2023.0.1.x86_64      docker-25.0.6-1.amzn2023.0.1.x86_64      iptables-libs-1
  iptables-nft-1.8.8-3.amzn2023.0.2.x86_64      libcgrou-3.0-1.amzn2023.0.1.x86_64      libnetfilter_cc
  libnftnl-1.0.1-19.amzn2023.0.2.x86_64      libnftnl-1.2.2-2.amzn2023.0.2.x86_64      pigz-2.5-1.amzn
  runc-1.1.11-1.amzn2023.0.1.x86_64

```

Complete!

```
[root@ip-172-31-90-172 ec2-user]# systemctl start docker
```

```
[root@ip-172-31-90-172 ec2-user]# sudo su
```

```
[root@ip-172-31-90-172 ec2-user]# yum repolist
```

```

repo id                                repo name
amazonlinux                            Amazon Linux 2023 repository

```

Run the command `yum repolist` to check whether the kubernetes repo has installed or not if successful installed then you can see a repo named as kubernetes

```

[root@ip-172-31-90-172 ec2-user]# yum repolist
repo id                                repo name
amazonlinux                            Amazon Linux 2023 repository
kernel-livepatch                       Amazon Linux 2023 Kernel Livepatch repository
kubernetes                             Kubernetes
[root@ip-172-31-90-172 ec2-user]# █

```

Do the above steps for all the instances i.e for worker-1 and worker-2.

5. Perform this ONLY on the Master machine. Initialize the Kubecuster

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all
```

```
[addons] Applied essential addon: kube-proxy
```

```
Your Kubernetes control-plane has initialized successfully!
```

```
To start using your cluster, you need to run the following as a regular user:
```

```

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

```

```
You should now deploy a pod network to the cluster.
```

```
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
```

```
https://kubernetes.io/docs/concepts/cluster-administration/addons/
```

```
Then you can join any number of worker nodes by running the following on each as root:
```

```

kubeadm join 172.31.12.28:6443 --token 4bqwb8.lua2ud01lr02uu55 \
--discovery-token-ca-cert-hash sha256:b4edc7948be9bca50767f623b58e0612feedc144a7364f95be8dbd8c4614a169

```

Copy the join command and keep it in a notepad, we'll need it later.

Copy the mkdir and chown commands from the top and execute them

```

[ec2-user@ip-172.31.12.28 docker]$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

```

Then, add a common networking plugin called flannel file as mentioned in the code.

```
kubectl apply -f
```

```
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

```
[ec2-user@ip-172.31.12.28 docker]$ kubectl apply -f
```

```
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

Check the created pod using this command

Now, keep a watch on all nodes using the following command - `watch kubectl get nodes`

6. Perform this ONLY on the worker machines

Run the following command

```
sudo yum install iproute-tc -y
```

```
sudo systemctl enable kubelet
```

```
sudo systemctl restart kubelet
```

Check the status of the pods using the following command

This command will show the status of all the pods.

```
kubectl get pods -n kube-system
```

Following command will show the status of the pod named daemonset.

```
kubectl get daemonset -n kube-system
```

```
[ec2-user@ip-172.31.12.28 docker]$ kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-55cb5b8774-fx12f	1/1	Running	0	100s
coredns-55cb5b8774-xn14v	1/1	Running	0	100s
etcd-ip-172.31.12.28.ec2.internal	1/1	Running	0	75s
kube-apiserver-ip-172.31.12.28.ec2.internal	1/1	Running	1	2m
kube-controller-manager-ip-172.31.12.28.ec2.internal	0/1	CrashLoopBackOff	1	70s
kube-proxy-4dv8m	1/1	Running	2	100s
kube-scheduler-ip-172.31.12.28.ec2.internal	1/1	Running	1	76s


```
[ec2-user@ip-172.31.12.28 docker]$ kubectl get daemonset -n kube-system
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
kube-proxy	1	1	1	1	1	kubernetes.io/os=linux	3m

That's it, we now have a Kubernetes cluster running across 3 AWS EC2 Instances. This cluster can be used to further deploy applications and their loads being distributed across these machines.

Conclusion:

Kubernetes cluster was successfully established using three AWS EC2 instances, which includes one Master and two Worker nodes. The process began with the creation of instances and configuration of settings to begin the communication. Docker was installed on all machines followed by the installation of Kubernetes components and the necessary repositories. The Master node was initialized with the `kubeadm init` command, and a plugin called Flannel was deployed to enable pod communication. Performing correct commands on the Worker nodes ensured they joined the cluster effectively. Also necessary commands confirmed the status of pods which indicated the proper working. Overall, the experiment provided information about working of the deployment and management of containerized applications in a distributed environment.